

Deep Learning

Lecture 8: Sequential Models

Amir Atapour-Abarghouei

amir.atapour-abarghouei@durham.ac.uk

Durham University





1 Recurrent neural networks

- Definition and implementation
- Backpropagation through time
- vanishing/exploding gradients

2 Long short-term memory

- Definition
- Properties
- Seq2Seq
- Attention

3 Transformers


- Self-Attention
- End-to-end object detection
- GPT
- DALL-E




Definition: recurrent neural networks

A function applied to nodes on a directed graph. *Unlike* one-way directed graphs (e.g. text, audio), sequential data is modelled using a cyclic connection that allows information to be stored [Rumelhart et al., 1986]. The same function f is applied to inputs at each time step, updating a hidden state vector h which acts as the network's memory:

$$h_{t+1} = f_{\theta}(h_t, x_t)$$




Computational Graphs




One-to-One

*Feedforward
network*




One-to-Many

e.g., image captioning



Many-to-One

e.g., sentence classification



Many-to-Many

seq2seq

Recurrent Neural Networks

Implementation




Example: RNN


A simple implementation is:

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

which is visually interpreted as a “cell”:



$$P(\mathbf{x}) = \prod_t P(x_t | x_{t-1}, x_{t-2}, \dots, x_1)$$



Recurrent Neural Networks

Backpropagation




Definition:

Backpropagation applied to an **unrolled** RNN is called backprop through time (BPTT). Gradients accumulate in W additively:

$$\frac{\partial \mathcal{L}_T}{\partial W} = \sum_{t \leq T} \frac{\partial \mathcal{L}_T}{\partial h_t} \frac{\partial h_t}{\partial W}$$

Long sequences use truncated BPTT where sequences are split into batches but hidden connections remain.



Let's Look at Some Code!




Simple RNN Example for Text Classification

Code on GitHub : https://github.com/atapour/dl-pytorch/blob/main/RNN_Sentiment_Analysis/RNN_Sentiment_Analysis.ipynb

Code on Colab: https://colab.research.google.com/github/atapour/dl-pytorch/blob/main/RNN_Sentiment_Analysis/RNN_Sentiment_Analysis.ipynb

Recurrent Neural Networks

Exploding and Vanishing Gradients




Why do gradients vanish/explode?

The gradient of involves many factors of W (and $\tanh$).

The product of T matrices converges to 0 (or grows to infinity) at an exponential rate in T .

$$\frac{\partial \mathcal{L}_T}{\partial W} = \sum_{t \leq T} \frac{\partial \mathcal{L}_T}{\partial h_t} \frac{\partial h_t}{\partial W} = \sum_{t \leq T} \frac{\partial \mathcal{L}_T}{\partial h_T} \frac{\partial h_T}{\partial h_t} \frac{\partial h_t}{\partial W}$$

Solution to exploding gradients: **clip gradients**



Long Short-Term Memory

Preventing Vanishing Gradient



Definition: long short term memory

LSTMs [Hochreiter et al., 1997] learn longer sequences than vanilla RNNs using better gradient flow. Backpropagation from c_t to c_{t-1} has no direct matrix multiplication by W .

Gates determine how much information passes through.

Always adding new information to the hidden state can be overwhelming.

Sometimes we want to forget things!

Example: LSTM cell



Gates control how much information is passed through using a sigmoid function and a dot product.



LSTM Properties

Main Strengths

- Allows for variable length sequences
- Efficient parameter usage
- Theoretically able to store arbitrarily old information

Main Limitations

- Practically unable to store very long term dependencies
- Limited by fixed size of hidden state
- Slow training and synthesis

Further Reading:

<https://karpathy.github.io/2015/05/21/rnn-effectiveness/>



- Seq2Seq (sequence to sequence) receives a sentence as the input and produces a sentence as the output.
- Seq2Seq will include two networks, one encoder and one decoder.
- An example of this would be “**Machine Translation**”:

Source Sentence (English)

granny liked my dishes.



Target Sentence (French)

mamie a aimé mes mets.




Seq2Seq Architecture

Machine Translation



- Let's see how inference works. We will take about the training later.

The sentence representation is used as the initial hidden state for the decoder.




Seq2Seq Architecture

Machine Translation



- Now, let's see how we would train this.

- The system is trained “end to end” as one single entity.



Seq2Seq Architecture


Issue



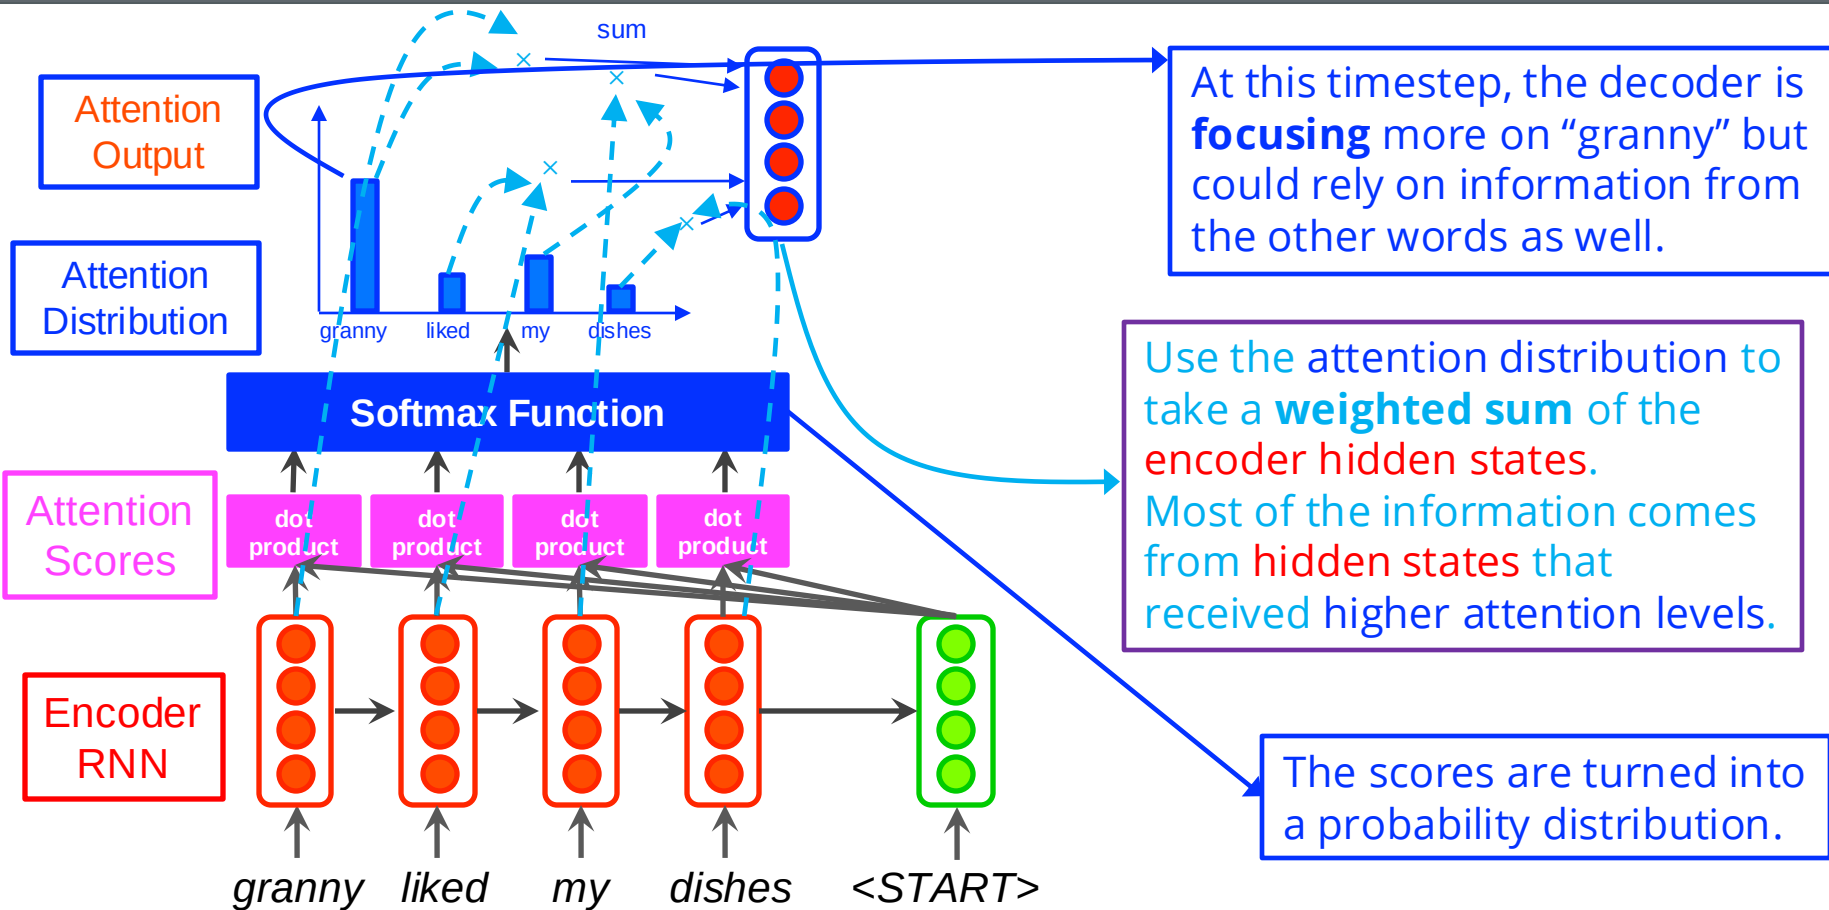
Solution: **Attention**

The entirety of the source sentence should be represented within the final hidden state of the encoder.


Informational Bottleneck:
too much pressure on a single vector




Seq2Seq with Attention



Seq2Seq with Attention




Seq2Seq with Attention



Decoder RNN

Seq2Seq with Attention





1. Encoder hidden states: $h_1, \dots, h_N \in \mathbb{R}^h$
2. At timestep t , decoder hidden state: $s_t \in \mathbb{R}^h$
3. Attention scores at timestep t : $e^t = [s_t^T h_1, \dots, s_t^T h_N] \in \mathbb{R}^N$
4. Softmax to get attention distribution: $\alpha^t = \text{softmax}(e^t) \in \mathbb{R}^N$
5. Use α^t to take a weighted sum of the encoder hidden states to get the attention output: $a_t = \sum_{i=1}^N \alpha_i^t h_i \in \mathbb{R}^h$
6. Concatenate attention output a_t with the decoder state s_t and continue with the rest of model training: $[a_t; s_t] \in \mathbb{R}^{2h}$

Encoder hidden states: $h_1, \dots, h_N$

Decoder hidden state: s_t

Definition: Given a set of vectors *values* and a vector *query*, attention is a weighted sum of the *values*, dependant on the *query*.


Attention types: ◦ Dot product ◦ Additive ◦ Multiplicative

Attention Is All You Need - Transformers



[Vaswani et al., 2017]

- **Self-Attention** provided a route to removing recurrence.
- **No RNNs.**




Attention Is All You Need - Transformers



[Vaswani et al., 2017]

- **Self-Attention** provided a route to removing recurrence.
- **No RNNs.**




Attention Is All You Need - Transformers



[Vaswani et al., 2017]

- **Self-Attention** provided a route to removing recurrence.
- **No RNNs.**




Attention Is All You Need - Transformers

[Vaswani et al., 2017]



Parallelisable

Transformers offer improvements in performance and efficiency.



In the decoder, we **mask out** words that the model is not supposed to have seen yet.



GPT Training and Evaluation

GPT-3 [Brown et al., 2020]:

- 175B parameters (96 layers with 96 heads each with 12,228 neurons)
- Batch size 3.2M. Input length of 2048
- Terabytes of data from the internet

GPT-4:

- **ChatGPT:** much larger (probably)

Evaluation Tasks

- Few shot translation
- Reading comprehension (Q&A)
- Closed book question & answering
- Natural language inference
- News article writing
- Dialogue generation

Example: ChatGPT article



ChatGPT

Title: United Methodists Agree to Historic Split

Subtitle: Those who oppose gay marriage will form their own denomination

The United Methodist Church, one of the largest Protestant denominations in the United States, has reached a pivotal moment in its history. After years of internal conflict and debate over LGBTQ+ rights, particularly the issue of gay marriage and the ordination of LGBTQ+ clergy, the denomination has agreed to a historic split. This decision will see the formation of a new, conservative denomination by members who oppose gay marriage.

The schism is the culmination of a long-standing rift within the church, reflecting broader societal

Alternative LLMs

- Google AI's Gemini
- Meta AI's LLaMA
- Anthropic's Claude
- Falcon
- ...




Example: Codex

Codex is GPT-3 trained on multiple datasets based on millions of GitHub repositories [one 2020 dataset consisted of 54 million public repos].

Available as an extension for Visual Studio Code and other IDEs.

VS Code Cursor + Sonnet 3.5

Output still contains bugs and errors but improving.



DALL-E Training

DALL-E [Ramesh et al., 2021]:

- 12-billion parameter version of GPT-3
- Generate images from detailed text descriptions
- **DALL-E 2** and now **DALL-E 3**


Capable of anthropomorphised versions of animals and objects, combining unrelated concepts in plausible ways, rendering text, and applying transformations to existing images

Example: DALL-E Image

A robot teaching his grandmother to suck eggs.



Good-looking man with golden hair and big bushy beard.







DEtection TRansformer

Fast object detection is crucial for many tasks including self-driving cars. Training end to end is difficult due to the discrete nature of objects.

DETR [Carion et al., 2020] enables global search and 'query' of the image for information. Attention matrices can also be used to make segmentation maps.

Example: architecture and examples



What we learned today!



1 Recurrent neural networks

- Definition and implementation
- Backpropagation through time
- vanishing/exploding gradients

2 Long short-term memory

- Definition
- Properties
- Seq2Seq
- Attention

3 Transformers

- Self-Attention
- End-to-end object detection
- GPT-3
- DALL-E